

Mining Massive Datasets: Week 1: MapReduce

Ian Quah (itq)

June 9, 2017

Topic 1

Map, Reduce and program structure

1. Map: (input elements) \rightarrow (0 ++ key-value pairs)
2. Reduce (key $[V_1, V_2, \dots, V_n]$) \rightarrow (0 ++ key-value pairs)
 - (a) If the reducer function is associative and commutative (order doesn't matter) can push off computation to map function.
 - (b) This is the concept behind a combiner: (map \rightarrow reduce) \rightarrow reduce
Examples:
 - i. sum is associative and commutative
 - ii. average is not, however, if we store the aggregate as (sum, # elems), we can reduce the number of tuples required to be transferred
3. Functions provided by the user:
 - (a) map
 - (b) reduce
 - (c) hash (optional): for user-specified grouping + specifying where to send jobs to
4. Structure of Program
 - (a) Master node sends data to all compute nodes, then periodically pings them
 - (b) Dealing with Failure:
 - i. Only failure on master can bring the whole thing down
 - ii. Bc master can know when things are down (ping), it can detect if a node failed if it did, the entire calculation must be restarted (as data for reduce was stored in failed node)
 - iii. Master assigns the original task to whichever cluster frees up, then notifies reduce that the source of incoming data is changing

Topic 2

Relational algebra

1. Definitions:

- Attributes: Column headers
- Relation: table with attributes
- tuples: rows of the relation
- Schema: set of attributes of a relation
- $R(A_1, A_2, \dots, A_n)$: - Relation name: R
- attributes: A_1, \dots, A_n

1. Relational Algebra Operations and definitions

- Selection:
 - Apply condition C to each tuple.
 - Returns only tuples that satisfy C
 - Denoted $\sigma_C(R)$
- Projection:
 - For some subset S of attributes of R
 - produce from each tuple only components for attributes in S
 - Denoted $\pi_S(R)$
- Union, Intersection, Difference:
 - Duh
- Natural Join:
 - Denoted as $R_1 \bowtie R_2$
 - Given R_1, R_2 , compare each pair of tuples
 - If agree on ALL attributes common between the two:
 - * return tuple s.t
 - * **Cond 1**: has components for each attribute in either schema
 - * **Cond 2**: agrees with two tuples on all attribs
 - else:
 - * return None
 - Discussion below talks about Nat Joins, but also applies to *equijoins* - equality covers equality of attributes that do not necessarily have the same name
- Grouping and Aggregation:
 - Duh